

Spam Detection using NLP — Summary & Insights

Project Objective

Build a machine learning model to accurately classify SMS messages as **spam** or **ham (not spam)** using Natural Language Processing (NLP) techniques, with careful consideration of **class imbalance**.

Phase 1: Prototyping & Experimentation (File1)

Initial experiments were conducted on a **smaller subset** of the dataset to efficiently evaluate various pipelines.

Key Challenge:

- The dataset is highly **imbalanced** (~87% ham vs ~13% spam)
- Using **accuracy** as a metric would be misleading, since predicting only "ham" yields high accuracy without real predictive power

Metric Selection:

The **F1-Score** was selected, especially for the minority spam class.

It balances **precision** (avoiding false positives) and **recall** (catching true spam), making it ideal for imbalanced classification.

Pipelines Compared:

Several text processing and model combinations were explored:

- Variations included **CountVectorizer**, **TF-IDF**, and different classifiers

Best Performer Identified:

- **TF-IDF Vectorizer + Logistic Regression**
-

Phase 2: Final Model on Full Dataset (File2)

After identifying the optimal pipeline, it was retrained on the **entire dataset** for full-scale evaluation and deployment.

Final Pipeline:

```
Pipeline([  
    ('tfidf', TfidfVectorizer()),  
    ('clf', LogisticRegression())
```

])

Evaluation Approach:

- **GridSearchCV** was used for hyperparameter tuning
- Evaluation was performed using `classification_report` with **Precision, Recall, and F1-score**

Typical results:

Class Precision Recall F1-score

Ham ~0.97 ~0.99 ~0.98

Spam ~0.95 ~0.89 ~0.92

Interpretation:

- High **precision for spam**: Few false alarms
- Strong **recall**: Most spam messages are detected
- The model performs reliably and balances both classes well

Final Step: Model Persistence

The trained pipeline was saved using `joblib`, making it ready for production or inference with minimal overhead.

Final Summary:

A practical, two-phase workflow was followed — beginning with prototype evaluation on a sample dataset and finalizing with full-scale model training and tuning. The resulting pipeline (TF-IDF + Logistic Regression) demonstrated strong performance, especially on the imbalanced spam class, and was saved for future use. This approach reflects best practices in model selection, evaluation, and deployment readiness.